


# Neural Networks







# Neural Networks

## Introduction

- Deep Learning - Rebirth of neural networks
- Inspired by the human brain (networks of neurons)




# History of Neural Networks



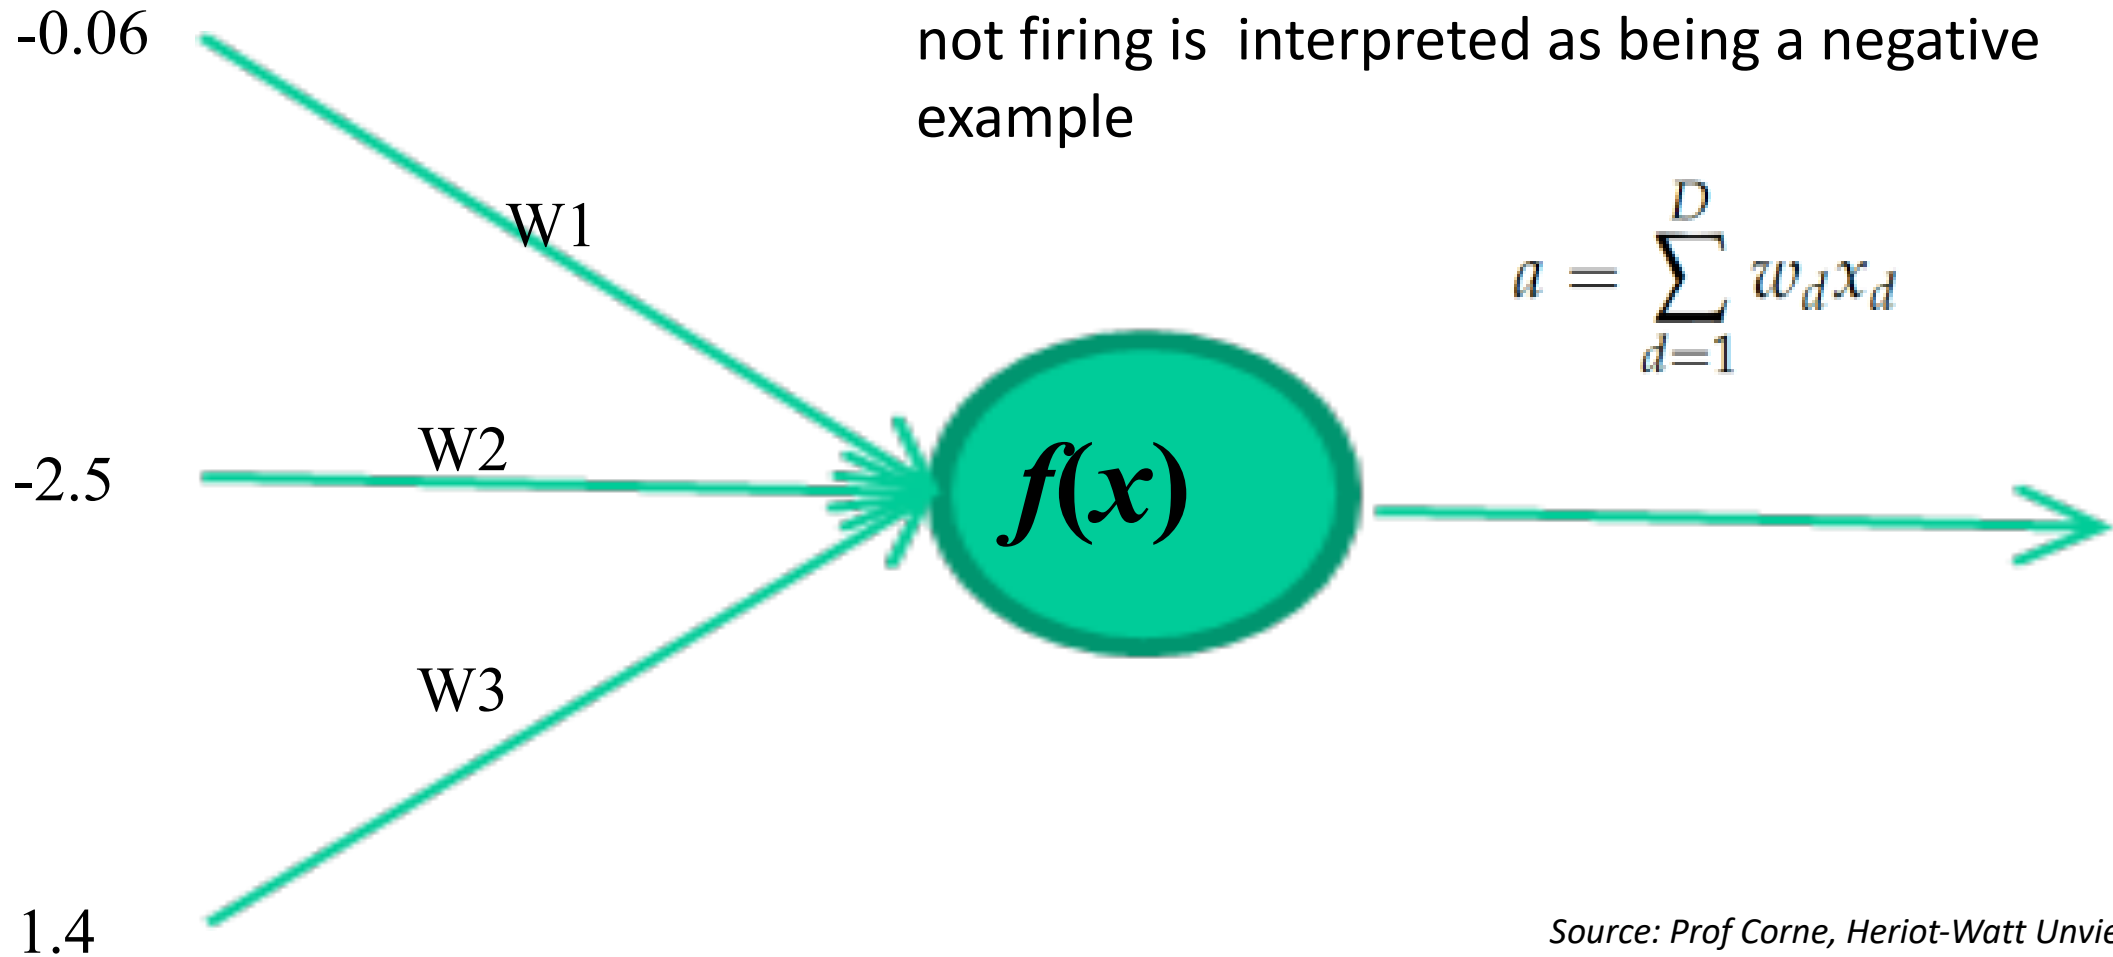
# Neuron

- Based on how much these incoming neurons are firing, and how “strong” the neural connections are, our main neuron will “decide” how strongly it wants to fire.
- Learning in the brain happens by neurons becoming connected to other neurons, and the strengths of connections adapting over time.
- Receives input from D-many other neurons, one for each input feature. The strength of these inputs are the feature values.




# Neuron as a Function

Firing is interpreted as being a positive example and not firing is interpreted as being a negative example



Source: Prof Corne, Heriot-Watt University, UK

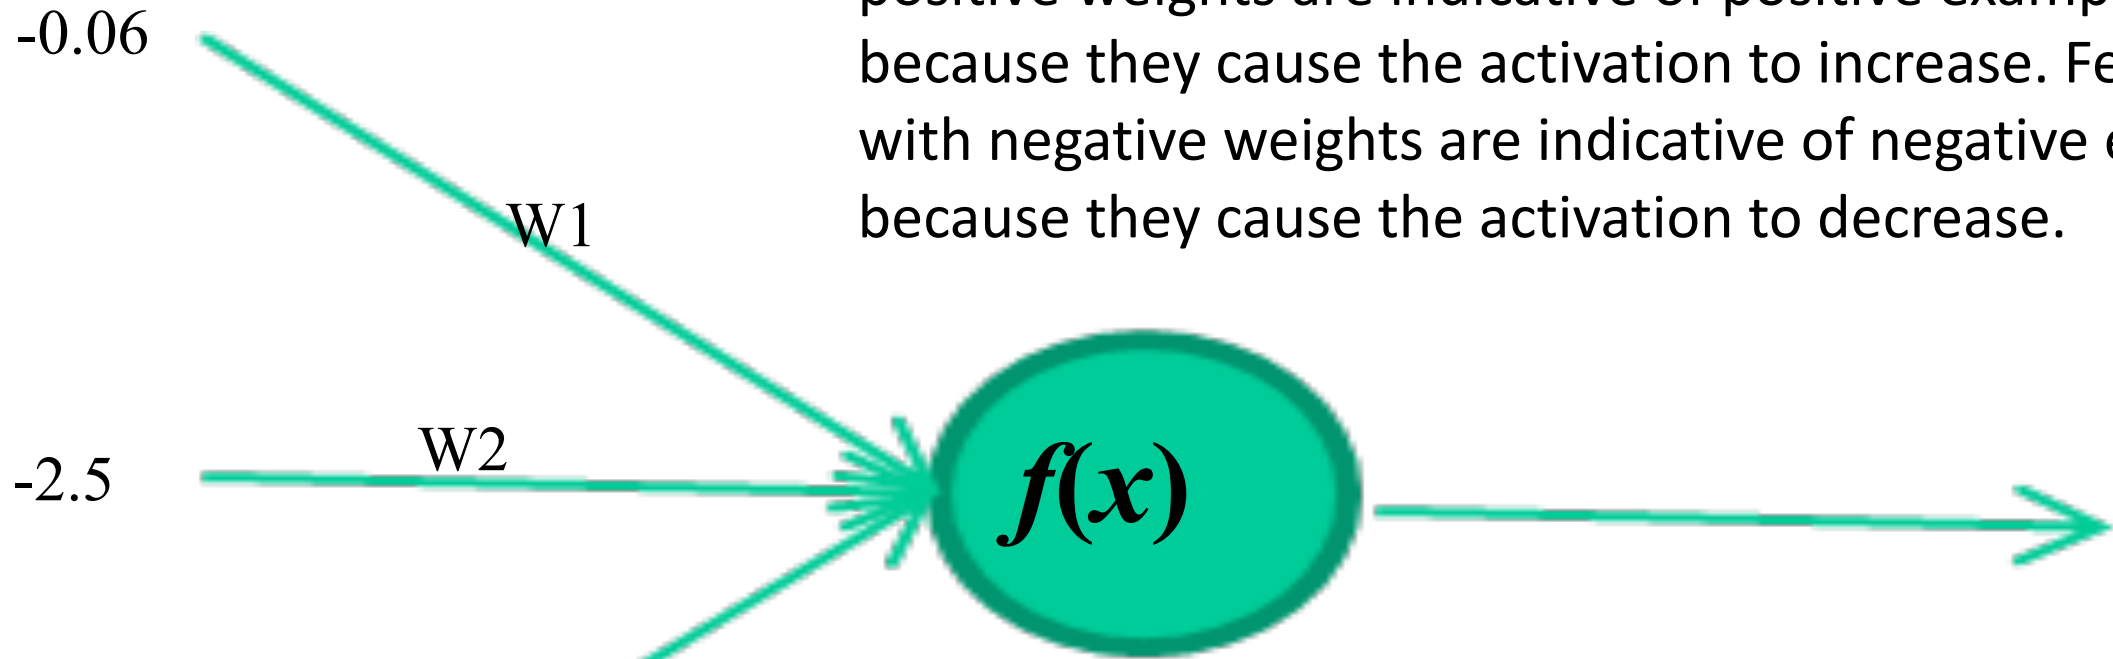
# Neuron as a Function



Source: Prof Corne, Heriot-Watt University, UK

# Neuron as a Function

Features with zero weight are ignored. Features with positive weights are indicative of positive examples because they cause the activation to increase. Features with negative weights are indicative of negative examples because they cause the activation to decrease.






$$x = -0.06 \times 2.7 + 2.5 \times 8.6 + 1.4 \times 0.002 = 21.34$$

$$a = \left[ \sum_{d=1}^D w_d x_d \right] + b$$

Source: Prof Corne, Heriot-Watt University, UK

# Perceptron Decision Boundary

$$\mathcal{B} = \left\{ x : \sum_d w_d x_d = 0 \right\}$$





# Perceptron Algorithm

---

**Algorithm 5** PERCEPTRONTRAIN( $\mathbf{D}$ ,  $MaxIter$ )

---

```
1:  $w_d \leftarrow 0$ , for all  $d = 1 \dots D$                                 // initialize weights
2:  $b \leftarrow 0$                                                     // initialize bias
3: for  $iter = 1 \dots MaxIter$  do
4:   for all  $(x, y) \in \mathbf{D}$  do
5:      $a \leftarrow \sum_{d=1}^D w_d x_d + b$                                 // compute activation for this example
6:     if  $ya \leq 0$  then
7:        $w_d \leftarrow w_d + yx_d$ , for all  $d = 1 \dots D$             // update weights
8:        $b \leftarrow b + y$                                            // update bias
9:     end if
10:  end for
11: end for
12: return  $w_0, w_1, \dots, w_D, b$ 
```

---

# Perceptron Algorithm

- Weight  $w_d$  is increased by  $yx_d$  and the bias is increased by  $y$ .
- Goal of the update is to adjust the parameters so that they are “better” for the current example



$$\begin{aligned} a' &= \sum_{d=1}^D w'_d x_d + b' \\ &= \sum_{d=1}^D (w_d + x_d) x_d + (b + 1) \\ &= \sum_{d=1}^D w_d x_d + b + \sum_{d=1}^D x_d x_d + 1 \\ &= a + \sum_{d=1}^D x_d^2 + 1 > a \end{aligned}$$

# Perceptron Algorithm

- **Online.** This means that instead of considering the entire data set at the same time, it only ever looks at one example.
- **Error driven.** This means that, so long as it is doing well, it doesn't bother updating its parameters

# Perceptron Algorithm

- If we make many many passes over the training data, then the algorithm is likely to overfit. On the other hand, going over the data only one time might lead to underfitting
- Loop over all the training examples in a constant
- Order is a bad idea. Re-permute the examples in each iteration.



# Perceptron Learning

- Iteratively pick a misclassified samples from the training set and apply the perceptron rule
- Each iteration through the training set is an *epoch*
- Continue training until total training set error ceases to improve (convergence)
- Perceptron Convergence Theorem: Guaranteed to find a solution in finite time if a solution exists
  - No. of required iterations  $\leq (R/\gamma)^2$

# Disadvantages of Perceptrons


- Can only learn linear decision boundaries
- Requires data to be linearly separable to guarantee convergence
- Even if data is linearly separable, convergence may take long
- Does not generalize easily to more than 2 classes
- Does not provide probabilistic outputs

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*A dataset*

<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

<i><b>Fields</b></i>	<i><b>class</b></i>
----------------------	---------------------

1.4 2.7 1.9	0
-------------	---


3.8 3.4 3.2	0
-------------	---

6.4 2.8 1.7	1
-------------	---

4.1 0.1 0.2	0
-------------	---

etc ...

Initialise with random weights



Source: Prof Corne, Heriot-Watt University, UK



# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Present a training pattern**



*Source: Prof Corne, Heriot-Watt University, UK*

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Compare with target output



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

<i>Fields</i>	<i>class</i>
---------------	--------------

1.4 2.7 1.9	0
-------------	---


3.8 3.4 3.2	0
-------------	---

6.4 2.8 1.7	1
-------------	---

4.1 0.1 0.2	0
-------------	---

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Present a training pattern**



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Compare with target output**



Source: Prof Corne, Heriot-Watt University, UK

# Neural Networks - Multi Layer Perceptrons

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                  0


3.8 3.4 3.2                  0

6.4 2.8 1.7                  1

4.1 0.1 0.2                  0

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK



# Neural Networks - Multi Layer Perceptrons

How do they learn?

Source: Prof Corne, Heriot-Watt University, UK

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0

3.8 3.4 3.2                      0


6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments, reduce the error


And so on ....



Called “Gradient Descent”


# The decision boundary perspective...

Initial random weights




# The decision boundary perspective...

**Present a training instance / adjust the weights**




# The decision boundary perspective...

**Present a training instance / adjust the weights**




# The decision boundary perspective...

**Present a training instance / adjust the weights**




# The decision boundary perspective...

**Present a training instance / adjust the weights**



# The decision boundary perspective...

Eventually ....



# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

First attempt at a training algorithm

- 1. **Initialize** network with **random** weights
- 2. **For all** training cases (**called examples**):
  - **a.** Present training inputs to network and calculate output
  - **b.** For all layers (starting with output layer, back to input layer):
    - i. Compare **network output** with **correct output** (error function)
    - ii. **Adapt weights** in current layer



# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Method for **learning weights** in feed-forward (FF) nets
- Can't use Perceptron Learning Rule
  - no **teacher values** are possible for **hidden units**
- Use **gradient descent** to minimize the error
  - **propagate deltas** to **adjust for errors**
  - **backward from outputs** to hidden layers **to inputs**

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

The idea of the algorithm can be summarized as follows :

1. Computes the **error term for the output units** using the observed error.
2. From output layer, repeat
  - propagating the error term back to the previous layer and **updating the weights between the two layers** until the earliest hidden layer is reached.

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Initialize weights (typically random!)
- In each epoch, do
  - **For each** example  $x^j$  in training set do
    - **forward pass** to compute
      - $y_{\text{pred}} = \text{NN}(x^j)$
      - error =  $(y^j - y_{\text{pred}})$  at each output unit
    - **backward pass** to calculate deltas to correct weights
    - update all weights
  - end
- Repeat until **training set error stops improving**

# Neural Networks Training: Backpropagation

## Gradient Descent

- Think of the  $N$  weights **as a point** in an  $N$ -dimensional space
- Add a **dimension** for the observed error
- Try to **minimize your position** on the “error surface”



# Neural Networks Training: Backpropagation

## Gradient Descent

- Trying to make **error decrease the fastest**
- **Compute:**
  - $\text{Grad}_E = [dE/dw_1, dE/dw_2, \dots, dE/dw_n]$
- **Change  $i^{\text{th}}$  weight by**
  - $\text{delta}_{w_i} = -\text{alpha} * dE/dw_i$
- We need a **derivative!**
- Activation function must be continuous, differentiable, non-decreasing, and easy to compute

# Neural Networks Training: Backpropagation

## Updating Hidden-to-Output

- We have **teacher supplied** desired values

- $$\text{delta}_{w_{ji}} = \alpha * (t_i - y_i) * g'(z_i) * a_j$$
$$= \alpha * (t_i - y_i) * y_i * (1 - y_i) * a_j$$

Here we have general formula with derivative, next we use for sigmoid

— for sigmoid the derivative is,  $g'(x) = g(x) * (1 - g(x))$

*Learning rate*

miss

Derivative of activation function

# Neural Networks Training: Backpropagation

## Updating interior weights

- Layer k units provide values to all layer k+1 units
  - “miss” is *sum of misses* from all units on k+1
  - $\text{miss}_j = \sum [ a_i(1 - a_i) (t_i - a_i) w_{ji} ]$
  - weights coming into this unit are *adjusted based on their contribution*

$$\text{delta}_{kj} = \alpha * I_k * a_j * (1 - a_j) * \text{miss}_j$$

Compute deltas

For layer k+1

# Making Choices


## Backpropagation

- Number of hidden layers – *empirically determined*
  - Too few ==> can't learn
  - Too many ==> poor generalization
- Number of neurons in each hidden layer – *empirically determined*
- Activation functions
- Error/loss functions
- Learning rate
- Gradient descent methods




# Activation Functions

Sigmoid




$$y = \frac{1}{1 + e^{-x}}$$

tanh




$$y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

ReLU



$$y = \max(0, x)$$

Leaky ReLU



$$y = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{if } \textit{otherwise} \end{cases}$$

# Loss Functions

- Euclidean loss / Squared loss  $L = \frac{1}{2} \|x_i - y_i\|_2^2$ 
  - Derivative w.r.t.  $x_i$   $\frac{\partial L}{\partial x_i} = x_i - y_i$
- Soft-max loss/multinomial logistic regression loss

$$p_i = \frac{e^{x_i}}{\sum_k e^{x_k}} \quad L = - \sum_i y_i \log(p_i)$$

- Derivative w.r.t.  $x_i$   $\frac{\partial L}{\partial x_i} = p_i - y_i$
- Also called: Cross-entropy loss

# Gradient Descent Methods

- Batch gradient descent (vs) Stochastic gradient descent (vs) Mini-batch stochastic gradient descent
  - Mini-batch SGD the most popularly used
- Using momentum
- Setting learning rate
  - Fixed learning rate
  - Using learning rate schedules
  - Adaptive learning rate methods: Adam, Adadelata, Adagrad, RMSProp

# Readings

- [“Introduction to Machine Learning” by Ethem Alpaydin](#), Chapters 11.1-11.11
- Bishop, PRML, Sec 5.1-5.3, 5.5
- Perceptron Convergence proof:  
<https://www.cse.iitb.ac.in/~shivaram/teaching/old/cs344+386-s2017/resources/classnote-1.pdf>